

FairHiringBias-Mitigation-NU

This project investigates bias in machine learning hiring decisions and applies fairness-aware techniques to ensure equitable treatment across gender groups. Developed for the Nile University AI Fairness Challenge, the project explores how model performance and fairness can coexist in real-world decision systems.

It aims to develop machine learning models that predict candidate hiring decisions based on multiple attributes while ensuring fairness across genders.

Table of Contents

- [Installation](#)
- [Dataset Overview](#)
- [Data Preprocessing](#)
- [Model Architecture & Performance](#)
 - [Model Architecture](#)
 - [Model Performance](#)
- [Fairness Analysis](#)
 - [Plots](#)
 - [Metrics](#)
- [Explainability](#)
- [Bias Mitigation Techniques](#)
 - [Reweighting](#)
 - [Feature Debiasing](#)

Installation

To install all the required dependencies, run:

```
pip install -r requirements.txt
```

Dataset Overview

The dataset contains **1,500 records with 11 columns**, representing various features relevant to hiring decisions. Each row corresponds to a unique job applicant, and the goal is to **predict whether the applicant was hired based on their attributes**.

- **Observations:**
 - No missing values in any column (the dataset is clean and complete).
 - About **50.8%** male and **49.2%** female data.
 - Data has no outliers.
 - The dataset includes a sensitive attribute **Gender**, which is the focus for fairness analysis.
 - Features are already numerically encoded, making them suitable for machine learning models.

Data Preprocessing

To ensure consistent scaling and appropriate handling of categorical variables, several preprocessing steps were applied:

- **Min-Max Scaling:**
 - Continuous numerical features such as `Age`, `EducationLevel`, `ExperienceYears`, `PreviousCompanies`, `DistanceFromCompany`, `InterviewScore`, `SkillScore`, and `PersonalityScore` were normalized using `MinMaxScaler` from `sklearn.preprocessing`. This transformation scales all values into the range `0, 1` improving the performance and convergence behavior of many ML algorithms.

The fitted scaler was also saved using `joblib` for reproducibility and potential use during inference (`minmax_scaler.pkl`).

```
joblib.dump(scaler, 'minmax_scaler.pkl')
```

- **One-Hot Encoding:**
 - The categorical feature `RecruitmentStrategy` was transformed into binary indicator columns using one-hot encoding via `pd.get_dummies()`.
 - This prevents the model from assuming any ordinal relationship between different strategies, which is crucial for unbiased learning.

These preprocessing steps resulted in a fully numerical and normalized dataset, ready for training fairness-aware machine learning models.

Model Architecture & Performance

To begin modeling, the dataset was first split into features (X) and the target variable `HiringDecision`.

- **A stratified train-test split** was then applied to ensure that the distribution of the hiring decision labels was preserved across both training and testing sets.
- This stratification is essential when dealing with **imbalanced classes** to ensure consistent evaluation.

```
# Sample 80% of males, 20% of females
biased_train = pd.concat([
    df_male.sample(frac=0.8, random_state=42),
    df_female.sample(frac=0.2, random_state=42)
])

# Shuffle the biased training set
biased_train = biased_train.sample(frac=1.0,
    random_state=42).reset_index(drop=True)
```

1. Model Architecture

1.1 Initial Model (Logistic Regression)

Logistic Regression classifier was trained using this biased dataset.

- The model was trained with a maximum of **1000 iterations** for convergence and a fixed random seed **random_state=42** for reproducibility.

1.2 Best Model (Random Forest)

Random Forest Classifier was trained with carefully tuned hyperparameters. This ensemble learning method aggregates the predictions of multiple decision trees, which enhances model robustness and generalization.

Model Hyperparameters:

```
rf_clf = RandomForestClassifier(  
    n_estimators=80,  
    max_depth=20,  
    min_samples_split=7,  
    min_samples_leaf=6,  
    max_features='sqrt',  
    random_state=42  
)
```

- This model achieved **the highest performance** before applying fairness mitigation techniques.

2. Model Performance

2.1 Logistic Regression Classifier Results

- **Training Accuracy:** 87.2%
- **Test Accuracy:** 85.33%

Class / Metric	Precision	Recall	F1-Score
Class 0	0.88	0.92	0.90
Class 1	0.80	0.71	0.75
Macro Avg	0.84	0.81	0.82
Weighted Avg	0.85	0.85	0.85
Accuracy	–	–	0.85

2.2 Random Forest Classifier Results

- **Training Accuracy:** 91.5%
- **Test Accuracy:** 88.7%

Class / Metric	Precision	Recall	F1-Score
Class 0	0.89	0.95	0.92
Class 1	0.86	0.75	0.80

Class / Metric	Precision	Recall	F1-Score
Macro Avg	0.88	0.85	0.86
Weighted Avg	0.89	0.89	0.88
Accuracy	–	–	0.8867

Fairness Analysis

To assess the fairness of the model, **Demographic Parity**, **Equal Opportunity**, and **Average Odds Difference** were evaluated with respect to the **Gender** attribute.

1. Plots

The plot shows the distribution of hiring predictions by **gender**:

 Gender Distribution

Gender	Predicted Not Hired (0)	Predicted Hired (1)
Male (0)	112	32
Female (1)	107	49

2. Metrics:

- **Demographic Parity:** **-0.0919**
- **Equal Opportunity:** **-0.1802**
- **Average Odds Difference:** **-0.0855**

Conclusion

- **Females were hired 9.2% more often than males overall.**
- Qualified females (y_true=1) had **18%** higher recall than qualified males.
- On average, **the model favors females by 8.6%.**

Explainability

To interpret the model's hiring decisions, **SHAP (SHapley Additive exPlanations)** was used with a **TreeExplainer on the trained Random Forest model**. The steps:

- **Sample Selection:** 3 candidates predicted as **Hired** and 2 as **Not Hired** were selected from the **test set**.
- **SHAP Computation:** SHAP values are computed for class 1 (**Hire**) for these 5 samples.
- **Force plots** were generated to show each feature's impact on the prediction:
 - **Red** = features pushing towards **Hire**.
 - **Blue** = features pushing towards **Not Hire**.

These plots helped confirm that **gender** had **near-zero influence**, while performance-related features (**InterviewScore**, **SkillScore**, **PersonalityScore**) were **the most influential**.

Explanation of Sample_3:

The applicant was predicted as **Hired** with **82% confidence**.

- These features increased the hiring score significantly:
 - **Strategy_3 = 0.0**
 - Not being recruited via Strategy 3 helped the candidate. (Maybe Strategy 3 is associated with weaker applicants in training.)
 - **ExperienceYears = 0.87**(normalized)
 - This is very high experience — a strong positive factor.
 - **SkillScore = 0.83**
 - A high skill score heavily contributing to being hired.
 - **Strategy_2 = 0.0**
 - Not being recruited via Strategy 2 was a slight boost.
 - **Strategy_1 = 1.0**
 - Being recruited via Strategy 1 helped.
- These features decreased the hiring probability, but not enough to outweigh the positives:
 - **EducationLevel = 2.0**
 - Mid-level education — perhaps candidates with higher education were generally favored in training.
 - **PersonalityScore = 0.35**
 - Below-average — had some negative influence.

Feature Importance

To show which features contribute most on average to predictions for class 1 (**Hire**) across the selected 100 samples (sample_X), a **global feature importance bar plot** was generated using **SHAP values for class 1 (Hire)**.

```
shap.summary_plot(shap_values_class1, sample_X, plot_type="bar")
```

Conclusion

- **Gender** has almost **zero SHAP importance** which means the model isn't directly using the Gender feature to make hiring decisions.

Earlier we found that:

- **Demographic Parity:** ~9.2% more likely to hire females.
- **Equal Opportunity:** ~18% more likely to correctly hire qualified females.

This suggests:

- Indirect bias is present.
- **Gender** may be correlated with other features (like lower scores or different strategy distributions).
- **Indirect Bias Flow:** **Gender** → **Strategy** → **Score Distributions** → **Prediction**.

SHAP results align well with correlation matrix, and both validate each other.

Bias Mitigation Techniques

To address gender-related bias in hiring decisions, multiple bias mitigation techniques were explored and evaluated their effects on model fairness and performance.

1. Reweighing

The Reweighing technique from the AIF360 library was applied to adjust instance weights in the training dataset based on group membership **Gender** and label **HiringDecision**. This method aims to reduce bias by assigning different weights to privileged (Male) and unprivileged (Female) groups to ensure fairness during training.

```
rw = Reweighing(
    privileged_groups=privileged_groups,
    unprivileged_groups=unprivileged_groups
)
train_rw = rw.fit_transform(train_bld)
```

- The Random Forest Classifier (best-performing model) was trained with the **same hyperparameters** as before on the reweighted dataset and evaluated its fairness.

However, the fairness metrics worsened slightly after reweighing compared to the original biased dataset.

2. Feature Debiasing

To improve fairness, Feature Debiasing approach was adopted manually by removing or minimizing the influence of the Gender feature.

- The Random Forest Classifier was trained again using the same parameters after applying this mitigation strategy.

Fairness Metrics After Feature Debiasing:

Metric	Value
Demographic Parity (FD)	-0.0849
Equal Opportunity (FD)	-0.2052
Average Odds Difference (FD)	-0.0884

Bias Mitigation Comparison

Metric	Before Mitigation	After Reweighing	After Feature Debiasing
Test Accuracy	0.8867	0.8833	0.8767
Demographic Parity	-0.0919	-0.0983	-0.0849

Metric	Before Mitigation	After Reweighing	After Feature Debiasing
Equal Opportunity	-0.1802	-0.1802	-0.2052
Average Odds Difference	-0.0855	-0.0904	-0.0884

Final Conclusion

- Equal Opportunity worsened after feature debiasing.

Despite removing Gender, the model still exhibits indirect bias (possibly due to proxy variables that correlate with gender (strategy, personality_score, interview_score)).